

# Rewards App Optimization

## Project Overview

Customer rewards programs are meant to increase customer spending and thus increase revenue. Bringing a customer into a business increases the likelihood of them spending more than the initial reward. Some customers only spend enough to get the reward and nothing more. Other customers spend regardless of a reward and may not even know there was a reward in the first place. If the customer didn't know there was a reward and they complete it then the reward is not the driving factor behind the revenue and could be considered a loss, because it didn't initiate any behavior.

The data set provided is fake generated data, but it helps to simulation real-world scenarios. This sample data can be used to build prediction models around which customer should not be offered a reward because it was not a behavior cause.

The sample data consists of customer profiles, e.g. income, gender, enrollment date; offers, e.g. discount/informational, duration, delivery method, etc.; and transactions, e.g. purchases, offers received/viewed/completed, rewards earned.

## Problem Statement

A method to reduce reward offers when unnecessary is needed. Unnecessary offers reduce company revenue which is not the intent. An offer should increase revenue by driving additional sales.

## Metrics

The output of the analysis is a binary answer to the question: Should this customer be sent a reward offer? The model results will be graded using both a confusion matrix, a ROC curve and a classification report.

The number of customers that shouldn't be given offers is very low compared to the number that should. This means that simply sending all customers an offer would result in a high accuracy score but does not solve the problem statement. This requires many different metrics to analyze various models for feasibility.

## Confusion Matrix

The confusion matrix shows the number of true positives, false positives, true negatives and false negatives. This chart is a good visual representation of what the model outputs.

## Recall (Sensitivity)

This value shows the true positive rate of the model. I want a model that has the best recall rate which should minimize false negatives.

## Analysis

### Data Preprocessing

The sample data contained three data sets: portfolio, profile, and transcript. Each data set went through a cleansing stage before being used in the main exploration stage.

#### Portfolio

This contains the offers that were made to customers. An offer can be informational, a discount, or buy-one-get-one (BOGO). Each offer has a duration and reward amount that would be given to a customer should it be completed within the timeframe.

As provided in the instruction set, informational offers do not have a reward but should be viewed as the impact of reading the information could impact purchase decisions for the duration indicated.

The data was enhanced as follows:

- Parse the delivery method (channel) into individual values. Raw data was in list format.
  - Email, mobile, social, web
- Add Boolean indicators for each offer type.
  - BOGO, informational, discount
- Convert duration from days to hours

While the delivery method (channel) of the offers could have valuable information, none of the offers used only a single method. Without individual delivery methods I was not able to decide if one method was better than another for certain customers.




#### Profile

This is the customer information; age, income, gender and join date. If this were actual customer data, it is likely provided by the customer and thus cannot be considered completely accurate. A customer can list any income and age or choose to provide nothing at all.

There were no duplicate customers, using the id provided, in the data. This meant no de-duplication was necessary.

If age was null than gender and income were also null. There were no cases where just one field was null. An age of 118 means no value was provided, i.e. null.

Below is the distribution of the raw data.



Cleaning the profile data was a bit more involved. I decided on these steps:

1. Add a new field called `age_null` that would be set to true if `age = 118`
2. Create age brackets
  - a. Under 18, 18-24, 25-34, 35-44, 45-54, 55-64, 65+
3. Add `income_null` set to True when income was null and set income to 0.
4. Set null gender values to 'Missing'.

### Transcript

This data set contains all the transactions a customer made and when. When they received offers, when they were viewed, when they were completed, and what rewards were earned. The events could contain multiple pieces of information. Specifically, if an offer was completed it's a single event but it

contains two values, which offer\_id and how much reward. I parsed this data to have one row per event, which turned reward into an event instead of having it listed under the offer\_completed event. This makes the data easier to work with.

I also standardized the event names since some had underscores, and others had spaces.

This is the largest dataset but required minimal transformation.

## Data Exploration

In order to answer the problem statement, I need to know, of the existing customers, which ones shouldn't be sent which offers. At it's simplest I need to find which customers were given offers and didn't view them but still completed the offer. This defines a set of customers that were going to make purchases regardless of an offer being extended.

The challenge is that I want this to be extensible to new customers, the cold start problem. New customers don't have any purchase history to build metrics with. This means that the input to a logistic regression model should only require the basic information that we get from the customer when they sign up for the program. As they make purchases and engage with the rewards program their behavior will change the models and determine whether they start receiving specific offers to drive behavior.

Using the historical transactions for each customer I can build this dataset. For all transactions I can aggregate the data to find offer\_recieved, offer\_viewed, and offer\_completed for every offer made to each customer. Using this base data, I can add a new column for offer\_unnecessary when a customer completes but does not view the offer. Now I know whether each offer a customer received was necessary.

A customer can also receive the same type of offer multiple times, i.e. customer receives more than one discount offer. This results in further aggregation by offer type which transforms the offer\_unnecessary from a Boolean to a probability. If the customer was given 2 offers and 1 of them was unnecessary then the offer\_unnecessary\_proba value is .5 or 50%. Notice that this does not apply to informational. These offers will always show 0% because they cannot be completed. Analysis on customer behavior is a different analysis that is outside the scope of this effort.

My final dataset looks like this, for example:

person	offer_type	offer_unnecessary_proba
68be06ca386d4c31939f3a4f0e3dd783	discount	0.0
0610b486422d4921ae7d2bf64640c50b	bogo	1.0
0610b486422d4921ae7d2bf64640c50b	informational	0.0

At this point a decision needs to be made. What is the cutoff for when an offer is deemed unnecessary? Is it 50%, 80% or something else? For this analysis I chose 80%.

For input into the models I encoded offer\_type and gender using ordinal encoding. Then excluded the original offer\_type and gender fields.

My model input dataset (X) is as follows:

	age	became_member_on	income	age_null	under_18	18-24	25-34	35-44	45-54	55-64	65+	income_null	offer_type_enc	gender_enc
0	118	20170212	0.0	True	False	False	False	False	False	False	False	True	1.0	2.0
1	55	20170715	112000.0	False	False	False	False	False	False	True	False	False	0.0	0.0
2	55	20170715	112000.0	False	False	False	False	False	False	True	False	False	2.0	0.0
4	75	20170509	100000.0	False	False	False	False	False	False	False	True	False	0.0	0.0
5	75	20170509	100000.0	False	False	False	False	False	False	False	True	False	2.0	0.0

## Results

### Model Evaluation and Validation

I split the above dataset into training and testing, 80% and 20%, respectively. The testing data will be held out to use on the final models.

I ran the training data through 5 different models using GridSearchCV to tune each one. I set GridSearchCV to use a custom scorer based on the recall score. This will rank the models by the recall score instead of the default score mechanism. I also used 5-fold cross-validation on the training dataset.

I initially included GradientBoostingClassifier but the training time was extremely long. If the model takes too long to train in a test environment, then it would probably not be a good fit for production if the model is retraining itself automatically. Chosen models must be performant.

Below is the outcome of the model training. Even though XGBClassifier has a higher overall score, I would choose RandomForestClassifier because it had the best score against the test data.

```

LogisticRegression
  Best Params      : {'class_weight': 'balanced', 'max_iter': 50, 'solver': 'liblinear'}
  Train Recall Score : 0.6062242983186863
  Test  Recall Score : 0.6302779963283505
  All   Recall Score : 0.6155778363165065

AdaBoostClassifier
  Best Params      : {'algorithm': 'SAMME', 'learning_rate': 0.5, 'n_estimators': 10}
  Train Recall Score : 0.5
  Test  Recall Score : 0.5
  All   Recall Score : 0.5

ExtraTreesClassifier
  Best Params      : {'class_weight': 'balanced_subsample', 'criterion': 'gini', 'max_depth': None, 'max_features': 0.5, 'max_leaf_nodes': 200, 'min_samples_leaf': 10, 'min_samples_split': 20, 'n_estimators': 100}
  Train Recall Score : 0.6969811032845936
  Test  Recall Score : 0.6908719231201529
  All   Recall Score : 0.7632489574277033

RandomForestClassifier
  Best Params      : {'class_weight': 'balanced_subsample', 'criterion': 'entropy', 'max_depth': None, 'max_features': 0.5, 'max_leaf_nodes': 100, 'min_samples_leaf': 30, 'min_samples_split': 2, 'n_estimators': 1000}
  Train Recall Score : 0.6953709230519051
  Test  Recall Score : 0.6876451800232288
  All   Recall Score : 0.7626594823957256

XGBClassifier
  Best Params      : {'gamma': 1, 'learning_rate': 1, 'max_delta_step': 0, 'max_depth': 15, 'min_child_weight': 6, 'n_estimators': 1000, 'subsample': 0.5}
  Train Recall Score : 0.5353904792343573
  Test  Recall Score : 0.533730939267918
  All   Recall Score : 0.8423894361276774

```

Expanding on the above model I also ran the data through a scaler to see if it would provide any improvements. Specifically, I chose MinMaxScaler, which translates the data to values between 0 and 1 based on the minimum and maximum value of each feature. When the data isn't on the same scale it can degrade predictive performance.

As you can see, there was an improvement in most of the models. You'll also see that by scaling the data the parameters used in the best model may have changed, e.g. LogisticRegression changed the max\_iter parameter from 50 to 25.

From these results I would change to the ExtraTreesClassifier model, since it has the best recall score on the test data, albeit only slightly.

```
LogisticRegression
  Best Params      : {'class_weight': 'balanced', 'max_iter': 25, 'solver': 'liblinear'}
  Train Recall Score : 0.6229733212708313
  Test  Recall Score : 0.6444822224719944
  All   Recall Score : 0.6338316957143055


AdaBoostClassifier
  Best Params      : {'algorithm': 'SAMME', 'learning_rate': 0.5, 'n_estimators': 10}
  Train Recall Score : 0.5
  Test  Recall Score : 0.5
  All   Recall Score : 0.5

ExtraTreesClassifier
  Best Params      : {'class_weight': 'balanced_subsample', 'criterion': 'entropy', 'max_depth': None, 'max_features':
0.5, 'max_leaf_nodes': 200, 'min_samples_leaf': 15, 'min_samples_split': 20, 'n_estimators': 100}
  Train Recall Score : 0.6971205703589944
  Test  Recall Score : 0.6957799445505976
  All   Recall Score : 0.7579042001306784

RandomForestClassifier
  Best Params      : {'class_weight': 'balanced_subsample', 'criterion': 'entropy', 'max_depth': None, 'max_features':
0.5, 'max_leaf_nodes': 100, 'min_samples_leaf': 20, 'min_samples_split': 2, 'n_estimators': 1000}
  Train Recall Score : 0.694420805721097
  Test  Recall Score : 0.6914222022404556
  All   Recall Score : 0.7676644353863105

XGBClassifier
  Best Params      : {'gamma': 1, 'learning_rate': 1, 'max_delta_step': 0, 'max_depth': 20, 'min_child_weight': 6, 'n_e
stimators': 1000, 'subsample': 0.5}
  Train Recall Score : 0.5306860819821114
  Test  Recall Score : 0.530677475535574
  All   Recall Score : 0.8741511034420949
```

Looking at these models another way, using the scaled model results, I can use a confusion matrix to visually see the ratio of predictions, i.e. true positive, false positive, true negative, and false negative. These are based on predictions across all the data for each model.



## Conclusion

### Reflection

Overall, using predictive modeling I was able to correctly identify whether a customer should be sent an offer based only on their initial profile about 70% of the time. This method would require retraining as more customers sign-up for the program and make transactions.

I started down the path of ROI based on false positives and false negatives but decided it was better having an initial starting point, as this analysis discusses. I found it very fascinating that with just a relatively few features I was able to expand out into all sorts of insights, most of which aren't included here for brevity.

## Improvement

There are several areas of improvement that can be made to this analysis.

First, additional models could be added to the analysis, including chaining multiple models together. Try using a deep learning model that can retrain itself as more data comes into the system.

Second, more model parameters could be evaluated. I was hardware limited in this analysis which meant a lot of time was spent waiting for the model tuning to finish.

Third, add metrics that show the monetary cost of incorrect predictions. It's obvious that sending out unnecessary offers costs the company money. But subtler, is the change to revenue if the customer would have never received an offer but should have. Offers tend to increase customer spending, which means if they didn't get an offer they should have received, its lost revenue. A model chosen by this method would heavily penalize incorrect predictions, which could change the model used.